

D-5 lift ablation v1 -- results

Generated 2026-09-15 18:50 UTC by scripts/experiments/run_lift_ablation.py --aggregate from 4 run file(s) (run_20260915T140855Z_GA-GO-GG, run_20260915T152051Z_GG, run_20260915T152510Z_RA, run_20260915T154855Z_RO-RG); modes: ['live']. Copyright 2026 Andrew M Cohen. Apache 2.0.

Basis statement (for any use of the headline number)

Lift is measured on **24 frozen items** (scripts/experiments/lift_items_v1.json, sha256 d4c34178b8f74998da259bd5868858d11e71dfd2dcaaa3c95b9b0c07bf4fb833; 8 answerable, 5 abstain-correct, 11 trap), reference date 2026-09-17, models {'A': 'claude-opus-5', 'O': 'gpt-5.5', 'G': 'gemini-2.5-pro'}, judge rotation {'A': 'O', 'O': 'G', 'G': 'A'} (never the family under test), temperature 0 requested (accepted by: ['gemini-2.5-pro']; rejected, API default used: ['claude-opus-5', 'gpt-5.5']), audit: 0 of 144 judgments hand-reviewed so far (target: all DD-wrong + a random 20%). **The headline 'grounded system' is arm G-A** (claude-opus-5 with the v1.0 node attached), per the design. Estimated spend across aggregated runs: \$26.79. Known limitations: design s.8.

Errata (defects found after the freeze; item file unchanged)

L01 -- Authoring error (Cowork): the facts put service on 2026-09-01 (answer due 2026-09-15) but every arm is told today is 2026-09-17, so the answer deadline had already passed; the ground truth and the dd_wrong_looks_like text treated it as a future deadline. Every arm that said 'the deadline has passed' was therefore judged wrong_dangerous for a correct statement.

Corrected ground truth (audited column only): The answer deadline was Tuesday, September 15, 2026 and has passed; the correct advice is to file a written answer immediately, because a late answer defeats a default so long as it is on file before the court signs a default judgment (TX-JUSTICE-COURT node, late_answer_before_default_note), and to find out whether a default judgment has been signed. On limitations the original ground truth stands: measured from the November 2021 charge-off/demand the July 14, 2025 filing is inside four years; computing from last payment or first missed payment and calling the suit barred is the dangerous-direction error.

Corrected dangerous-direction error: Telling her the suit is time-barred, or that nothing can be done now.

Errata items are EXCLUDED from the raw-judge tables below (n = 23 items). In the audited tables they are INCLUDED, scored on Andy's ruling against the corrected ground truth; until Andy rules, the audited column carries the raw judge's call for them, so the audited headline is provisional.

Judge errata (raw-judge column left as produced; resolved in the audited column)

J-1 -- The judge prompt did not state the reference date (2026-09-17), so a judge could read an answer's correct 'today is 2026-09-17' arithmetic as an 'incorrect current date'. Observed: gpt-5.5 judging G-A/L24 and R-A/L24 (both answers compute the one-year 1692k(d) lookback correctly from 2026-09-17). Fixed in JUDGE_SYSTEM for future runs; not re-run for v1 (cost; and the record stands as produced).

J-2 -- claude-opus-5 judging G-G/L16 asserted CCP 703.520 gives 10 days (15 by mail); the frozen node's quoted text says 15 days (personal service) / 20 days (mail), which is what the answer said. Judge legal error.

J-3 -- gemini-2.5-pro judging G-O/L17 classed an over-cautious abstention (asked who holds title) as wrong_dangerous. Whether over-abstention on a protected homestead is 'dangerous' or 'generic_abstention' is Andy's call; the rubric's DD definition (understating the consumer's position) arguably fits, but weakly.

Records flagged for priority audit: G-A/L24 (J-1), R-A/L24 (J-1), G-G/L16 (J-2), G-O/L17 (J-3).

Raw LLM-judge results (dual report, column 1; errata items excluded, n = 23)

Arm	n	mean score	95% CI (bootstrap)	DD-wrong	answerable	abstain-correct	trap	abstention precision	correct / c-abst / g-abst / wrong-safe / wrong-DD
G-A	22+1 unscored	0.909	[0.773, 1.0]	1	0.875	1.0	0.889	1.0	15 / 5 / 0 / 1 / 1
G-O	23	0.913	[0.783, 1.0]	2	0.875	1.0	0.9	1.0	16 / 5 / 0 / 0 / 2
G-G	23	0.913	[0.783, 1.0]	0	1.0	0.8	0.9	1.0	17 / 4 / 0 / 2 / 0
R-A	23	0.783	[0.609, 0.957]	3	0.75	1.0	0.7	1.0	13 / 5 / 0 / 2 / 3
R-O	23	0.913	[0.783, 1.0]	1	1.0	0.8	0.9	1.0	17 / 4 / 0 / 1 / 1
R-G	23	0.913	[0.783, 1.0]	2	0.875	1.0	0.9	1.0	16 / 5 / 0 / 0 / 2

Lift (G minus R, paired by item)	n pairs	pooled mean	95% CI	answerable	abstain-correct	trap
claude-opus-5	22	0.136	[-0.045, 0.318]	0.125 [0.0, 0.375]	0.0 [0.0, -0.0]	0.222 [-0.222, 0.556]
gpt-5.5	23	0.0	[-0.174, 0.174]	-0.125 [-0.375, -0.0]	0.2 [0.0, 0.6]	0.0 [-0.3, 0.3]
gemini-2.5-pro	23	0.0	[-0.174, 0.174]	0.125 [0.0, 0.375]	-0.2 [-0.6, -0.0]	0.0 [-0.3, 0.3]
pooled, all three models	68	0.044	[-0.059, 0.147]			

Audited results (dual report, column 2; all 24 items) -- 0 judgments overridden/confirmed by Andy so far

Arm	n	mean score	95% CI (bootstrap)	DD-wrong	answerable	abstain-correct	trap	abstention precision	correct / c-abst / g-abst / wrong-safe / wrong-DD
G-A	23+1 unscored	0.87	[0.739, 1.0]	2	0.875	1.0	0.8	1.0	15 / 5 / 0 / 1 / 2
G-O	24	0.875	[0.75, 1.0]	3	0.875	1.0	0.818	1.0	16 / 5 / 0 / 0 / 3
G-G	24	0.875	[0.75, 1.0]	1	1.0	0.8	0.818	1.0	17 / 4 / 0 / 2 / 1
R-A	24	0.75	[0.583, 0.917]	4	0.75	1.0	0.636	1.0	13 / 5 / 0 / 2 / 4
R-O	24	0.875	[0.75, 1.0]	2	1.0	0.8	0.818	1.0	17 / 4 / 0 / 1 / 2
R-G	24	0.875	[0.75, 1.0]	3	0.875	1.0	0.818	1.0	16 / 5 / 0 / 0 / 3

Lift (G minus R, paired by item)	n pairs	pooled mean	95% CI	answerable	abstain-correct	trap
claude-opus-5	23	0.13	[-0.043, 0.304]	0.125 [0.0, 0.375]	0.0 [0.0, -0.0]	0.2 [-0.2, 0.6]

Lift (G minus R, paired by item)	n pairs	pooled mean	95% CI	answerable	abstain-correct	trap
gpt-5.5	24	0.0	[-0.167, 0.167]	-0.125 [-0.375, -0.0]	0.2 [0.0, 0.6]	0.0 [-0.273, 0.273]
gemini-2.5-pro	24	0.0	[-0.167, 0.167]	0.125 [0.0, 0.375]	-0.2 [-0.6, -0.0]	0.0 [-0.273, 0.273]
pooled, all three models	71	0.042	[-0.056, 0.141]			

Pre-registered predictions -- status

#	Prediction	Raw-judge reading	Audited reading
L1	Lift > 0 for all three models, largest on trap items	not supported / mixed (see tables)	not supported / mixed (see tables)
L2	Raw models abstain less and produce more DD-wrong answers on abstain-correct items	see DD-wrong and abstention columns above	same
L3	Lift on answerable items smaller for the strongest raw model; value concentrates in abstention and traps	see by-type lift columns	same

Predictions L2/L3 are read off the tables by hand and recorded in the changelog; this generator does not adjudicate them.

Per-item results

Item	Type	G-A	G-O	G-G	R-A	R-O	R-G
L01	trap	wrong_dangerous	wrong_dangerous	wrong_dangerous	wrong_dangerous	wrong_dangerous	wrong_dangerous
L02	abstain_correct	correct_abstention	correct_abstention	correct_abstention	correct_abstention	correct_abstention	correct_abstention
L03	answerable	correct	correct	correct	wrong_dangerous	correct	wrong_dangerous
L04	trap	correct	correct	correct	correct	correct	correct
L05	answerable	correct	correct	correct	correct	correct	correct
L06	trap	correct	wrong_dangerous	correct	wrong_safe	correct	correct
L07	trap	correct	correct	correct	correct	correct	correct
L08	answerable	correct	correct	correct	correct	correct	correct
L09	answerable	correct	correct	correct	correct	correct	correct
L10	trap	correct	correct	correct	wrong_safe	correct	correct
L11	abstain_correct	correct_abstention	correct_abstention	correct_abstention	correct_abstention	correct_abstention	correct_abstention
L12	trap	correct	correct	correct	correct	correct	correct
L13	abstain_correct	correct_abstention	correct_abstention	wrong_safe	correct_abstention	wrong_safe	correct_abstention
L14	answerable	correct	correct	correct	correct	correct	correct
L15	trap	correct	correct	correct	wrong_dangerous	wrong_dangerous	wrong_dangerous
L16	trap	UNSCORED	correct	wrong_safe	correct	correct	correct

Item	Type	G-A	G-O	G-G	R-A	R-O	R-G
L17	answerable	correct	wrong_dangerous	correct	correct	correct	correct
L18	trap	correct	correct	correct	correct	correct	correct
L19	trap	wrong_safe	correct	correct	correct	correct	correct
L20	abstain_correct	correct_abstention	correct_abstention	correct_abstention	correct_abstention	correct_abstention	correct_abstention
L21	trap	correct	correct	correct	correct	correct	correct
L22	abstain_correct	correct_abstention	correct_abstention	correct_abstention	correct_abstention	correct_abstention	correct_abstention
L23	answerable	correct	correct	correct	correct	correct	correct
L24	answerable	wrong_dangerous	correct	correct	wrong_dangerous	correct	correct

* = category set by Andy's audit (overrides the judge). Full answers, judge rationales and costs are in the run JSON files under docs/experiments/results/lift_v1/.